

1 RAM Model

1.1 Memory

Infinite sequence of cells, contains w bits. Every cell has an address starting at 1

1.2 CPU

32 registers of width w bits.

1.2.1 Operations

Set value to register (constant or from other register). Take two integers from other registers and store the result of; $a+b$, $a-b$, $a \cdot b$, a/b . Take two registers and compare them; $a < b$, $a = b$, $a > b$. Read and write from memory.

1.3 Definitions

An algorithm is a set of atomic operations. It's cost is the number of atomic operations. A word is a sequence of w bits

2 Worst-case

Worst-case cost of an algorithm is the longest possible running time of input size n

3 Dictionary search

```
let  $n$  be register 1, and  $v$  be register 2
register  $left \rightarrow 1$ ,  $right \rightarrow 1$ 
while  $left \leq right$ 
    register  $mid \rightarrow (left + right)/2$ 
    if the memory cell at address  $mid = v$ 
    then
        return yes
    else if memory cell at address  $mid > v$ 
    then
         $right = mid - 1$ 
    else
         $left = mid + 1$ 
return no
```

Worst-case time: $f_2(n) = 2 + 6 \log_2 n$

4 Big-O

We say that $f(n)$ grows asymptotically no faster than $g(n)$ if there is a constant $c_1 > 0$ such that $f(n) \leq c_1 \cdot g(n)$ and holds for all n at least a constant c_2 . This is denoted by $f(n) = O(g(n))$.

$\lim_{n \rightarrow \infty} \frac{f(n)}{g(n)} = c$ for some constant c

4.1 Example

$1000 \log_2 n = O(n)$,
 $n \neq O(10000 \log_2 n)$
 $\log_{b_1} n = O(\log_{b_2} n)$ for any constants $b_1 > 1$ and $b_2 > 1$. Therefore $f(n) = 2 + 6 \log_2 n$ can be represented; $f(n) = O(\log n)$

5 Big-Ω

If $g(n) = O(f(n))$, then $f(n) = \Omega(g(n))$ to indicate that $f(n)$ grows asymptotically no slower than $g(n)$. We say that $f(n)$ grows asymptotically no slower than $g(n)$ if $c_1 > 0$ such $f(n) \geq c_1 \cdot g(n)$ for $n > c_2$; denoted by $f(n) = \Omega(g(n))$

6 Big-Θ

If $f(n) = O(g(n))$ and $f(n) = \Omega(g(n))$, then $f(n) = \Theta(g(n))$ to indicate that $f(n)$ grows asymptotically as fast as $g(n)$

7 Sort

7.1 Merge Sort

Divide the array into two parts, sort the individual arrays then combine the arrays together. $f(n) = O(n \log n)$.

This is the fastest sorting time possible (apart from $O(n \log \log n)$)

7.2 Counting Sort

A set S of n integers and every integer is in the range $[1, U]$. (all integers are distinct)

Step 1: Let A be the array storing S . Create array B of length U . Set B to zero.

Step 2: For $i \in [1, n]$; Set x to $A[i]$, Set $B[x] = 1$

Step 3: Clear A , For $x \in [1, U]$; If $B[x] = 0$ continue, otherwise append x to A

7.2.1 Analysis

Step 1 and 3 take $O(U)$ time, while Step 2 $O(n)$ time. Therefore running time is $O(n + U) = O(U)$.

8 Random

RANDOM(x, y) returns an integer between x and y chosen uniformly at random

9 Data

9.1 Data Structure

Data Structure describes how data is stored in memory.

9.2 LinkedList

Every node stores pointers to its succeeding and preceding nodes (if they exist). The first node is called the head and last called the tail. The space required for a linkedlist is $O(n)$ memory cells. Starting at the head node, the time to enumerate over all the integers is $O(n)$. Time for assertion and deletion is equal to $O(1)$

9.3 Stack

The stack has two operations; Push (Inserts a new element into the stack), Pop (Removes the most recently inserted element from the stack and returns it. Since a stack is just a linkedlist, push and pop use $O(1)$ time.

9.4 Queue

The queue has two operations; En-queue (Inserts a new element into the queue), De-queue (Removes the least recently used element from the queue and returns it). Since a queue is just a linkedlist, push and pop use $O(1)$ time.

10 Dynamic Arrays

10.1 Naive Algorithm

insert(e): Increase n by 1, initial an array A' of length n , copy all $n-1$ of A to A' , Set $A'[n]=e$, Destroy A .

This takes $O(n^2)$ time to do n insertions.

10.2 A Better Algorithm

insert(e): Append e to A and increase n by 1. If A is full; Create A' of length $2n$, Copy A to A' , Destroy A and replace with A' . This takes $O(n)$ time to do n insertions.

11 Hashing

The main idea of hashing is to divide the dataset S into a number m of disjoint subsets such that only one subset needs to be searched to answer any query.

11.1 Pre-processing

Create an array of linkedlist(L) from 1 to m and an array H of length m . Store the heads of L in H , for all $x \in S$; calculate hash value ($h(x)$), insert x into $L_{h(x)}$. We will always choose $m = O(n)$, so $O(n + m) = O(n)$

11.2 Querying

Query with value v , calculate the hash value $h(v)$, Look for v in $L_{h(v)}$. Query time: $O(|L_{h(v)}|)$

11.3 Hash Function

Pick a prime p ; $p \geq m$, $p \geq$ any integer k . Choose α and β uniformly random from $1, \dots, p-1$. Therefore: $h(k) = 1 + (((\alpha k + \beta) \bmod p) \bmod m)$

11.3.1 Any Possible Integer

The possible integers is finite under the RAM Model. Max: $2^w - 1$. Therefore p exists between $[2^w, w^{w+1}]$.

11.4 Timing

Space: $O(n)$, Preprocessing time: $O(n)$, Query time: $O(1)$ in expectation

12 Week 3 - Extra

When using 'Direction 1: Constant Finding' setting c_1 , always set it to match the coefficient on the LHS so that you can cancel.

When trying to get a contradiction, try and isolate an $x \cdot c_1$ on the RHS, where $x \in \mathbb{Z}$, such that an expression that contains n is $\leq x \cdot c_1$

Make judicious use of the \max function when adding functions together. If $f_1(n) + f_2(n) \leq c_1 \cdot g_1(n) + c'_1 \cdot g_2(n) \leq \max\{c_1, c'_1\} \cdot (g_1(n) + g_2(n))$, for all $n \geq \max\{c_2, c'_2\}$.

13 The Master Theorem

13.1 Theorem 1

$$n + \frac{n}{c} + \frac{n}{c^2} + \dots + \frac{n}{c^h} = O(n)$$

13.2 Theorem 2

Let $f(n)$ be a function that returns a positive value for every integer $n > 0$. We know:

$$f(1) \leq c_1$$

$$f(n) \leq \alpha \cdot f(\lceil n/\beta \rceil) + c_2 \cdot n^\gamma \text{ for } n \geq 2$$

where $\alpha, \beta, \gamma, c_1$ and c_2 are positive constants. Then:

- If $\log_b \alpha < \gamma$ then $f(n) = O(n^\gamma)$
- If $\log_b \alpha = \gamma$ then $f(n) = O(n^\gamma \cdot \log(n))$
- If $\log_b \alpha > \gamma$ then $f(n) = O(n^{\log_b \alpha})$

14 SE Set 3

Find out how many times a recurrence takes to terminate, and then proceed to eyeball the time complexity

15 Hierarchy

$$\begin{aligned} O(1) &\leq O(\log(n)) \leq O(n^c) \\ &\leq O(n) \leq O(n^2) \\ &\leq O(n^c) \leq O(c^n) \end{aligned}$$

16 Trees

An undirected graph is a pair of (V, E) where:

- V is a set of elements, each of which is called a node
- E is a set of pairs u, v such that:
 - u and v are distinct nodes;
 - If (u, v) is in E , then (v, u) is also in E - we say that there is an edge between u and v .

A node may also be called a vertex. We will refer to V as the vertex set or the node set of the graph, and E the edge set.

For example, there is a graph (V, E) where $V = \{a, b, c, d, e\}$ $E = \{(a, b), (b, a), (b, c), (c, b), (a, d), (d, a), (b, d), (d, b), (c, e), (e, c)\}$

The number of edges equals $|E|/2 = 10/2 = 5$

Let $G = (V, E)$ be an undirected graph. A path in G is a sequence of nodes (v_1, v_2, \dots, v_k) such that

- For every $i \in [1, k-1]$, there is an edge between v_i and v_{i+1}

A cycle in G is a path (v_1, v_2, \dots, v_k) such that

- $k \geq 4$

- $v_1 = v_k$
- v_1, v_2, \dots, v_k are distinct

An undirected graph $G = (V, E)$ is connected if, for any two distinct vertices u and v , G has a path from u to v .

A tree is connected undirected graph contains no cycles.

A tree with n nodes has $n - 1$ edges.

Given any tree T and an arbitrary node r , we can allocate a level to each node as follows:

- r is the root of T - this is level 0 of the tree
- All the nodes that are n edge away from r constitute level n of the tree

The number of levels is called the height of the tree. We say that T has been rooted once a root has been designated.

Consider a tree T that has been rooted.

Let u and v be two nodes in T . We say that u is the parent of v if:

- v is at the level immediately below u
- There is an edge below u and v

Accordingly, we say that v is a child of u .

Let u and v be two nodes in T . We say that u is the ancestor of v if one of the following holds:

- $v = u$
- u is the parent of v
- u is the parent of an ancestor of v

Accordingly, we say that v is a descendant of u .

In particular, if $u \neq v$, we say that u is a proper ancestor of v , and likewise, v is a proper descendant of u .

Let u be a node in a rooted tree T . The subtree of u is the part of T that is "at or below" u .

In a rooted tree, a node is a leaf node if it has no children; otherwise it is an internal node.

Let T be a rooted tree where every internal node has at least 2 child nodes. If m is the number of leaf nodes, then the number of internal nodes is at least $m - 1$.

A k -ary tree is a rooted tree where every internal node has at most k child nodes.

A 2-ary tree is called a binary tree.

Consider a binary tree with height h . Its level n ($0 \leq n \leq h - 1$) is full if it contains 2^n .

A binary tree of height h is complete if:

- Levels $0, 1, \dots, h - 2$ are all full
- At Level $h - 1$, the leaf nodes are "as far left as possible"

A complete tree with $n \geq 2$ nodes has height $O(\log n)$

